FLO STUDIOS · AI EXPLORATION · 2026

The Smarter the Tool, the More Expensive Your Vagueness Gets

Why AI amplifies unclear thinking instead of fixing it

Flo Studios · 6 min read · AI · Creator Economy · Software Development

There is an assumption quietly embedded in how most teams talk about AI, and it goes something like this: the tool is powerful enough to compensate for whatever we bring to it. If the output isn't good, we need a better model. If the response feels generic, we need to prompt it differently. If the result misses the brief, the brief probably wasn't detailed enough — but the model will figure it out. It's an assumption that feels reasonable until you watch it play out in practice, at which point it becomes one of the most expensive beliefs a team can hold.

AI doesn't compensate for unclear thinking. It scales it. A vague brief given to a human produces one mediocre output that a single person spent time on. A vague brief given to AI produces fifty mediocre outputs in four minutes, and now someone has to review all of them, form an opinion on all of them, and explain to the tool why none of them are right — without being able to articulate what right actually looks like. The cost of not knowing what you want has never been higher, and most teams are paying it without realising that's what the invoice is for.

> *AI doesn't compensate for unclear thinking. It scales it. A vague brief given to a human produces one mediocre output. A vague brief given to AI produces fifty — and now someone has to review all of them.*

What vagueness actually costs

To understand why this matters so much right now, it helps to understand what AI is actually doing when it receives an unclear instruction — because the model doesn't experience vagueness the way a person does. A person given a vague brief will ask a clarifying question, make an assumption and flag it, or apply judgment based on context they've accumulated over years of similar work. An LLM given a vague brief will produce something confident, structurally sound, and calibrated to the statistical average of everything it has ever seen that resembles this kind of request — which means it will give you the most generic possible version of what you asked for, delivered with complete conviction.

That's not a model failure. That's the model doing exactly what it was designed to do with the information it was given. The failure lives upstream, at the point where the person using the tool hadn't done the thinking that would have made the instruction specific. And here is what makes this so costly in a world where AI is embedded into every part of the creative and development process: the volume of output means the thinking deficit compounds at speed. Every unclear input doesn't produce one poor output that gets revised — it produces a cascade of outputs that each carry the shape of the original vagueness, and the team spends its energy reacting to the outputs rather than resolving the thinking that produced them.

The two kinds of vagueness

Not all unclear thinking looks the same, and the distinction matters because each kind has a different cost and a different fix.

The first is directional vagueness — not knowing what you actually want the output to do, who it's for, or what a successful result looks like in concrete terms. This is the most common kind and the most expensive, because no amount of iteration closes the gap when the destination hasn't been defined. A team that doesn't know whether a piece of content is meant to build trust, drive a specific action, or establish authority will produce ten drafts that each optimise for something slightly different, and the feedback loop will keep cycling because the evaluation criteria keep shifting.

The second is standard vagueness — knowing what you want in general terms but not having a sharp enough definition of quality to evaluate what comes back. This shows up as the feeling that the output is almost right but something's off, without being able to name what. It produces longer revision cycles, inconsistent results across a team using the same tool, and a growing frustration that gets attributed to the model rather than to the absence of a defined standard. The model isn't inconsistent. The team's picture of what good looks like is.

> *The model isn't inconsistent. The team's picture of what good looks like is — and AI makes that gap impossible to hide.*

Why AI makes this harder to ignore

Before AI was embedded in creative and development workflows, vagueness had a natural speed limit. The time it took a human to produce an output created space for the thinking to catch up. A brief would get written, someone would ask a question that exposed a gap, the gap would get filled, and the work would proceed with a slightly clearer picture of what it was trying to be. The process was slow enough that clarity could emerge through iteration without too much waste.

AI removed that speed limit entirely, which means the gap between thinking and doing has collapsed in a way that exposes every piece of unclear thinking immediately and at scale. There is no longer a natural checkpoint where vagueness gets surfaced before it becomes output — the output arrives before the thinking is finished, and the team is left managing the consequences of having moved faster than their clarity allowed. This is why so many teams feel like AI is making them busier rather than more productive: they are producing more, reviewing more, revising more, and making more decisions per hour than they ever have, but the proportion of that work that actually advances the project is lower than it should be because the inputs that generated it weren't sharp enough.

What working clearly with AI actually looks like

The teams that get consistently strong output from AI tools are almost never the ones with the most sophisticated prompting techniques. They are the ones who do their thinking before they open the tool, and carry that thinking into every input they give it. In practice, this looks like a few specific habits that are less about technical knowledge and more about the discipline to answer hard questions before asking the model anything.

It means being able to state, in a sentence, what the output needs to do and for whom — not what it should contain, but what job it is supposed to perform and for what specific person. It means having a concrete example of output that is close to what you're looking for, even if it's imperfect, because examples communicate standard in a way that descriptions rarely do. It means knowing in advance what would make you reject a response — what specific quality, tone, or structural characteristic the output cannot have — because negative constraints are often clearer than positive ones and the model responds to both equally well.

And it means treating the first output not as a draft to be revised but as a diagnostic — a signal about the quality of the input rather than a starting point to be worked forward from. When the first response is off, the correct next step is almost always to sharpen the input, not to revise the output. Teams that understand this stop accumulating revision cycles and start accumulating clarity, which is the asset that compounds.

> *When the first response is off, the correct next step is to sharpen the input — not revise the output. One habit produces clarity. The other produces cycles.*

The thinking is the work

There is a version of AI adoption that treats the tool as a way to skip the hard part of creative and technical work — the part where you have to figure out what you actually want and why, where you have to develop a sharp enough standard to evaluate what you produce, where you have to make judgment calls about quality that can't be automated. That version produces a lot of output and very little progress, and the frustration it generates tends to land on the model.

The version that compounds is the one where the tool is used to accelerate thinking that is already clear — where the input carries the weight of the judgment that went into it, and the output is a function of that judgment rather than a substitute for it. In that version, AI genuinely extends what a team can do, because it removes the friction from execution without removing the human from the decision. The thinking remains the work. AI just means you can act on it faster once you've done it.

The cost of vagueness has always been real. AI just made it impossible to defer.